

```
import numpy as np
```

```
class Softmax(object):
```

```
    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)
```

```
    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001
```

```
    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.
```

Inputs:

- X: A numpy array of shape (N, D) containing a minibatch of data.
- y: A numpy array of shape (N,) containing training labels; y[i] = c means that X[i] has label c, where $0 \leq c < C$.

Returns a tuple of:

- loss as single float

```
    """
```

```
    # Initialize the loss to zero.
    loss = 0.0
```

```
    unit_loss = np.zeros(X.shape[0])
    # ===== #
    # YOUR CODE HERE:
    #     Calculate the normalized softmax loss. Store it as the variable loss.
    #     (That is, calculate the sum of the losses of all the training
    #     set margins, and then normalize the loss by the number of
    #     training examples.)
    # ===== #
    assert X.shape[1] == self.W.shape[1], f"{X.shape[1]} != {self.W.shape[1]}"
    def softmax(c, x):
        return np.exp(self.W[c]@x) / (np.exp(self.W@x).sum())
```

```
#     i = 0
    for y_i, x_i in zip(y, X):
        loss -= np.log(softmax(y_i, x_i))
    #     unit_loss[i] = -np.log(softmax(y_i, x_i))
    #     i += 1
```

```
    loss /= X.shape[0]
    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```
#     return loss, unit_loss
    return loss
```

```
def loss_and_grad(self, X, y):
    """
```

Same as `self.loss(X, y)`, except that it also returns the gradient.

Output: `grad` -- a matrix of the same dimensions as `W` containing the gradient of the loss with respect to `W`.

"""

Initialize the loss and gradient to zero.

`loss = 0.0`

`grad = np.zeros_like(self.W)`

=====

YOUR CODE HERE:

Calculate the softmax loss and the gradient. Store the gradient as the variable grad.

=====

assert `X.shape[1] == self.W.shape[1]`, `f"{X.shape[1]} != {self.W.shape[1]}"`

def softmax(`c, x`):

return `np.exp(self.W[c]@x) / (np.exp(self.W@x).sum())`

for `y_i, x_i` **in** `zip(y, X)`:

`loss -= np.log(softmax(y_i, x_i))`

`loss /= X.shape[0]`

for `i` **in** `range(0, X.shape[0])`:

for `k` **in** `range(0, self.W.shape[0])`:

`dL_i_dw_k = (softmax(k, X[i]) - (k == y[i]))*X[i]`

assert `grad[k,:].shape == dL_i_dw_k.shape`, `f"{grad[k,:].shape} != {dL_i_dw_k.sha`

`pe} =}"`

`grad[k,:] += dL_i_dw_k`

`grad /= X.shape[0]`

=====

END YOUR CODE HERE

=====

return `loss, grad`

def grad_check_sparse(`self, X, y, your_grad, num_checks=10, h=1e-5`):

"""

sample a few random elements and only return numerical in these dimensions.

"""

for `i` **in** `np.arange(num_checks)`:

`ix = tuple([np.random.randint(m) for m in self.W.shape])`

`oldval = self.W[ix]`

`self.W[ix] = oldval + h` *# increment by h*

`fxph = self.loss(X, y)`

`self.W[ix] = oldval - h` *# decrement by h*

`fxmh = self.loss(X,y)` *# evaluate f(x - h)*

`self.W[ix] = oldval` *# reset*

`grad_numerical = (fxph - fxmh) / (2 * h)`

`grad_analytic = your_grad[ix]`

`rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))`

print(`'numerical: %f analytic: %f, relative error: %e'` % (`grad_numerical, grad_analytic, rel_error`))

```

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #

    assert X.shape[1] == self.W.shape[1], f"{X.shape[1]} != {self.W.shape[1]}"
    loss = (-np.log((np.exp((self.W[y]*X).sum(axis=1))/(np.exp(X@self.W.T).sum(axis=1))))).mean()

    softmax_mat = ((np.exp(X@self.W.T))/(np.exp(X@self.W.T).sum(axis=1)[:,np.newaxis]))
    indicator_func = 1*(y[:,np.newaxis] == np.arange(10))
    grad = (1.0/X.shape[0])*(softmax_mat - indicator_func).T@X

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (batch_size, dim)

```

```
# - y_batch should have shape: (batch_size,)
# The indices should be randomly generated to reduce correlations
# in the dataset. Use np.random.choice. It's okay to sample with
# replacement.
# ===== #
idxs = np.random.choice(np.arange(num_train), size=batch_size, replace=False)
X_batch = X[idxs]
y_batch = y[idxs]
# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
# Update the parameters, self.W, with a gradient step
# ===== #
self.W -= learning_rate*grad

# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    # Predict the labels given the training data.
    # ===== #
    y_pred = np.argmax(X@self.W.T, axis=1)
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred
```