

```
import numpy as np
```

```
from .layers import *
from .layer_utils import *
```

```
class TwoLayerNet(object):
    """
```

A two-layer fully-connected neural network with ReLU nonlinearity and softmax loss that uses a modular layer design. We assume an input dimension of D, a hidden dimension of H, and perform classification over C classes.

The architecture should be affine - relu - affine - softmax.

Note that this class does not implement gradient descent; instead, it will interact with a separate Solver object that is responsible for running optimization.

The learnable parameters of the model are stored in the dictionary self.params that maps parameter names to numpy arrays.

```
    """
```

```
def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
              dropout=0, weight_scale=1e-3, reg=0.0):
```

```
    """
```

Initialize a new network.

Inputs:

- input\_dim: An integer giving the size of the input
  - hidden\_dims: An integer giving the size of the hidden layer
  - num\_classes: An integer giving the number of classes to classify
  - dropout: Scalar between 0 and 1 giving dropout strength.
  - weight\_scale: Scalar giving the standard deviation for random initialization of the weights.
  - reg: Scalar giving L2 regularization strength.
- ```
    """
```

```
self.params = {}
```

```
self.reg = reg
```

```
# ===== #
```

```
# YOUR CODE HERE:
```

```
# Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
# self.params['W2'], self.params['b1'] and self.params['b2']. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
# The dimensions of W1 should be (input_dim, hidden_dim) and the
# dimensions of W2 should be (hidden_dims, num_classes)
```

```
# ===== #
```

```
self.params['W1'] = weight_scale * np.random.randn(hidden_dims, input_dim)
self.params['b1'] = np.zeros(hidden_dims)
self.params['W2'] = weight_scale * np.random.randn(num_classes, hidden_dims)
self.params['b2'] = np.zeros(num_classes)
```

```
# ===== #
```

```
# END YOUR CODE HERE
```

```
# ===== #
```

```
def loss(self, X, y=None):
```

```
    """
```

Compute loss and gradient for a minibatch of data.

Inputs:

- X: Array of input data of shape (N, d<sub>1</sub>, ..., d<sub>k</sub>)
- y: Array of labels, of shape (N,). y[i] gives the label for X[i].

Returns:

If y is None, then run a test-time forward pass of the model and return:

- scores: Array of shape (N, C) giving classification scores, where scores[i, c] is the classification score for X[i] and class c.

If y is not None, then run a training-time forward and backward pass and return a tuple of:

- loss: Scalar value giving the loss
- grads: Dictionary with the same keys as self.params, mapping parameter names to gradients of the loss with respect to those parameters.

"""

scores = None

```
# ===== #
# YOUR CODE HERE:
#   Implement the forward pass of the two-layer neural network. Store
#   the class scores as the variable 'scores'. Be sure to use the layers
#   you prior implemented.
```

```
# ===== #
W1, b1 = self.params['W1'], self.params['b1']
W2, b2 = self.params['W2'], self.params['b2']
reg = self.reg
h1, cache_h1 = affine_relu_forward(X, W1, b1)
z2, cache_z2 = affine_forward(h1, W2, b2)
scores = z2
```

```
# ===== #
# END YOUR CODE HERE
# ===== #
```

# If y is None then we are in test mode so just return scores

```
if y is None:
    return scores
```

loss, grads = 0, {}

```
# ===== #
# YOUR CODE HERE:
#   Implement the backward pass of the two-layer neural net. Store
#   the loss as the variable 'loss' and store the gradients in the
#   'grads' dictionary. For the grads dictionary, grads['W1'] holds
#   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
#   i.e., grads[k] holds the gradient for self.params[k].
#
#   Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
#   for each W. Be sure to include the 0.5 multiplying factor to
#   match our implementation.
#
#   And be sure to use the layers you prior implemented.
```

```
# ===== #
loss, dL = softmax_loss(scores, y)
dh1, dw2, db2 = affine_backward(dL, cache_z2)
dx, dw1, db1 = affine_relu_backward(dh1, cache_h1)
```

```
reg_loss = reg*0.5*(np.linalg.norm(W1, ord='fro')**2 + np.linalg.norm(W2, ord='fro')**2)
loss += reg_loss
dw1 += reg*W1
dw2 += reg*W2
```

```
grads['W1'] = dw1
grads['b1'] = db1
```

```

grads['W2'] = dw2
grads['b2'] = db2

# ===== #
# END YOUR CODE HERE
# ===== #
return loss, grads

class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                  dropout=0, use_batchnorm=False, reg=0.0,
                  weight_scale=1e-2, dtype=np.float32, seed=None):
        """
        Initialize a new FullyConnectedNet.

        Inputs:
        - hidden_dims: A list of integers giving the size of each hidden layer.
        - input_dim: An integer giving the size of the input.
        - num_classes: An integer giving the number of classes to classify.
        - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
        the network should not use dropout at all.
        - use_batchnorm: Whether or not the network should use batch normalization.
        - reg: Scalar giving L2 regularization strength.
        - weight_scale: Scalar giving the standard deviation for random
        initialization of the weights.
        - dtype: A numpy datatype object; all computations will be performed using
        this datatype. float32 is faster but less accurate, so you should use
        float64 for numeric gradient checking.
        - seed: If not None, then pass this random seed to the dropout layers. This
        will make the dropout layers deterministic so we can gradient check the
        model.
        """
        self.use_batchnorm = use_batchnorm
        self.use_dropout = dropout > 0
        self.reg = reg
        self.num_layers = 1 + len(hidden_dims)
        self.dtype = dtype
        self.params = {}

        # ===== #
        # YOUR CODE HERE:
        # Initialize all parameters of the network in the self.params dictionary.
        # The weights and biases of layer 1 are W1 and b1; and in general the
        # weights and biases of layer i are Wi and bi. The
        # biases are initialized to zero and the weights are initialized
        # so that each parameter has mean 0 and standard deviation weight_scale.
        # ===== #

```

```

for i, dims in enumerate(hidden_dims):
    if self.use_batchnorm:
        prev_size = (input_dim if i == 0 else hidden_dims[i-1])
        self.params[f'gamma_{i+1}'] = np.ones(prev_size)
        self.params[f'beta_{i+1}'] = np.zeros(prev_size)
    if i == 0:
        self.params['W1'] = weight_scale * np.random.randn(hidden_dims[i], input_dim)
        self.params['b1'] = np.zeros(hidden_dims[i])
    else:
        self.params[f'W{i+1}'] = weight_scale * np.random.randn(hidden_dims[i], hidden_dims[i-1])
        self.params[f'b{i+1}'] = np.zeros(hidden_dims[i])
L = len(hidden_dims)+1
self.params[f'W{L}'] = weight_scale * np.random.randn(num_classes, hidden_dims[len(hidden_dims)-1])
self.params[f'b{L}'] = np.zeros(num_classes)

# ===== #
# END YOUR CODE HERE
# ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
if seed is not None:
    self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

```

If `y` is not `None`, then run a training-time forward and backward pass and return a tuple of:

- `loss`: Scalar value giving the loss
- `grads`: Dictionary with the same keys as `self.params`, mapping parameter names to gradients of the loss with respect to those parameters.

```
"""
```

```
X = X.astype(self.dtype)
```

```
mode = 'test' if y is None else 'train'
```

```
# Set train/test mode for batchnorm params and dropout param since they
# behave differently during training and testing.
```

```
if self.dropout_param is not None:
    self.dropout_param['mode'] = mode
```

```
if self.use_batchnorm:
    for bn_param in self.bn_params:
        bn_param[mode] = mode
```

```
scores = None
```

```
# ===== #
# YOUR CODE HERE:
```

```
# Implement the forward pass of the FC net and store the output
# scores as the variable "scores".
```

```
# ===== #
```

```
e = 1e-8
```

```
reg = self.reg
```

```
cache = []
```

```
state = np.reshape(X, (len(X), -1))
```

```
N = len(X)
```

```
for i in range(1, self.num_layers):
```

```
    assert len(state) == N
```

```
    # Assuming state: (N x D)
```

```
    state, cache_ = affine_forward(state, self.params[f'W{i}'].T, self.params[f'b{i}'])
```

```
    cache.append(cache_)
```

```
    assert len(state) == N
```

```
    # Assuming state: (N x M)
```

```
    if self.use_batchnorm:
```

```
        mu = state.mean(axis=0)[:, np.newaxis]
```

```
        var = state.var(axis=0)[:, np.newaxis]
```

```
        x_hat = (state - mu) / np.sqrt(var + e)
```

```
        cache.append((mu, var, x_hat, state))
```

```
        state = self.params[f'gamma_{i}'] * x_hat + self.params[f'beta_{i}']
```

```
    assert len(state) == N
```

```
state, cache_ = relu_forward(state)
```

```
cache.append(cache_)
```

```
assert len(state) == N
```

```
if self.use_dropout:
```

```
    if mode == 'test':
```

```
        state *= (1 - self.dropout_param['p'])
```

```
    else:
```

```
        M = np.random.rand((state.shape[1])) < self.dropout_param['p']
```

```
        state *= M
```

```
        cache.append(M)
```

```
    assert len(state) == N
```

```
L = self.num_layers
```

```
state, cache_ = affine_forward(state, self.params[f'W{L}'].T, self.params[f'b{L}'])
```

```
cache.append(cache_)
```

```
scores = state
```

```

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
# ===== #

loss, dL = softmax_loss(scores, y)
dL, dw, db = affine_backward(dL, cache.pop())
grads[f'W{L}'] = dw.T + reg*self.params[f'W{L}']
grads[f'b{L}'] = db
assert grads[f'W{L}'].shape == self.params[f'W{L}'].shape, f"{grads[f'W{L}'].shape=} | {self.params[f'W{L}'].shape=}"
assert grads[f'b{L}'].shape == self.params[f'b{L}'].shape, f"{grads[f'b{L}'].shape=} | {self.params[f'b{L}'].shape=}"
reg_loss = reg*0.5*(np.linalg.norm(self.params[f'W{L}'], ord='fro')**2)

for i in range(self.num_layers-1, 0, -1):
    if self.use_dropout:
        if mode == 'test':
            dL *= (1-self.dropout_param['p'])
        else:
            M = cache.pop()
            dL *= M

    dL = relu_backward(dL, cache.pop())

    # Assuming dL: (N x D) - Each row is a feature, each column an observation. Sum along
    # observations
    if self.use_batchnorm:
        mu, var, x_hat, x = cache.pop()
        dbeta = dL.sum(axis=0)
        dgamma = (dL*x_hat).sum(x=0)
        dx_hat = dL*self.params[f'gamma_{i}']
        dmu = (-1/(var+e))*dx_hat.sum(axis=0)
        dvar = (-0.5/((var+e)**(3.0/2.0))) * ((x - mu)*dx_hat).sum(axis=0)
        dL = 1/((var+e)**(1.0/2.0))*dx_hat + 2*((x - mu)/len(m))*dvar + (1.0/len(x))*dmu
        grads[f'beta_{i}'] = dbeta
        grads[f'gamma_{i}'] = dgamma
        assert grads[f'gamma_{i}'].shape == self.params[f'gamma_{i}'].shape, \
            f"{grads[f'gamma_{i}'].shape=} | {self.params[f'gamma_{i}'].shape=}"
        assert grads[f'beta_{i}'].shape == self.params[f'beta_{i}'].shape, \
            f"{grads[f'beta_{i}'].shape=} | {self.params[f'beta_{i}'].shape=}"

    dL, dw, db = affine_backward(dL, cache.pop())
    grads[f'W{i}'] = dw.T + reg*self.params[f'W{i}']
    grads[f'b{i}'] = db
    assert grads[f'W{i}'].shape == self.params[f'W{i}'].shape, f"{grads[f'W{i}'].shape=} | {self.params[f'W{i}'].shape=}"
    assert grads[f'b{i}'].shape == self.params[f'b{i}'].shape, f"{grads[f'b{i}'].shape=} | {self.params[f'b{i}'].shape=}"
    reg_loss += reg*0.5*(np.linalg.norm(self.params[f'W{i}'], ord='fro')**2)

loss += reg_loss

```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #  
return loss, grads
```